



## Frontier Engineering Challenge 2026

Hosted By - micro1

### CHALLENGE

Aug 28 – Aug 31, 2026

3:00 PM – 6:00 PM · UTC

🌐 Online · 👥 Team size: 1 · 👤 7.5K registrations



🔗 Help

[Overview](#) [About micro1](#) [Theme](#) [Instruction](#) [Prizes & Awards](#) [Stages & Timeline](#)

## Overview

**Build at the frontier of agentic AI. August 28–31, 2026. Online, individual, free.**

AI can produce convincing code in seconds. Real engineering begins when convincing is not enough: incomplete requirements, hidden dependencies, difficult edge cases, failure modes, and decisions that require technical judgment.

The micro1 Frontier Engineering Challenge is a free, global, online competition for engineers ready to work at that frontier. Over a three-day sprint you will use coding agents to tackle a real-world software engineering problem and produce a solution that is correct, reproducible, testable and clearly explained.

**\$10,000**

Prize pool



**Up to 50**

Paid opportunities at micro1



**Three**

Selective awards



[View Options](#) ^

# Process

1

## Register

Free, takes two minutes. Registration closes **Aug 29 at 23:59 UTC**.

2

## Build

Problem released at kickoff. Three days, coding agents required.

3

## Prove

Show a baseline and an improved solution, with evidence.

4

## Submit

Code, reproduction guide, video and agent trajectories.

The full challenge statement is released at kickoff (Aug 28 at 15:00 UTC). Until then, expect to build at the frontier of agentic AI without details that would reveal the problem in advance.

Participants will be able to register after the Hackathon kicks off. The Hackathon ends on Aug 31 at 18:00 UTC.

Participants are evaluated on the quality of the submitted work, not résumé prestige, job title, university, employer or geography.

**Coding-agent use is required.** You must disclose the tools you used and submit the required trajectories for evaluation.

---

## About micro1

micro1 is the leading AI data lab for training frontier models and evaluating AI agents. Experts contribute their diverse subject matter knowledge across domains such as finance, healthcare, STEM engineering, and more. micro1 transforms that real-world expertise into high-quality training data, evaluations, and feedback loops that improve how AI systems learn, reason, and perform.

Our platform identifies and vets top talent through an AI recruiter, enabling high-quality expert contributions at scale. We aim to enable 1 billion people to do meaningful work by applying their expertise to AI. As our global expert network grows, micro1 is building the human intelligence layer for frontier AI.

0:00 / 1:13

---

Stay connected with Micro1 across our social channels for the latest updates, opportunities, and insights: [LinkedIn](#) | [Instagram](#) | [X](#) | [Reddit](#) | [YouTube](#)

---


## Theme

### **Build at the frontier of agentic AI.**

Use coding agents to tackle a real-world engineering problem where correctness, reproducibility and human judgment matter.

In short: let your creativity run. Show us how you think through code. Apply it to any industry you like.

## Example use cases



The full problem PDF, starter materials, constraints and acceptance tests are published at kickoff.

**Every valid entry must present both a baseline solution and an advanced solution.** The advanced solution should show a meaningful improvement in capability, reliability, efficiency, coverage or engineering quality, not a cosmetic variation.

### Technology policy

Supported languages include Python, TypeScript, Java, C++, Go and Rust. Commonly used frameworks and libraries in those ecosystems are allowed, provided the entry stays reproducible and complies with the final problem PDF.

Illustrative examples: FastAPI, Flask, Django, LangGraph and related Python ecosystems; Node.js, Express, NestJS and Next.js for TypeScript; Spring Boot for Java; standard C++ and CMake toolchains; Go modules and common Go web frameworks; and Cargo, Tokio, Axum or Actix for Rust. This list is intentionally non-exhaustive.

The problem PDF may prescribe a starter repository, runtime, dependency limits, API access or testing environment where needed for fair and deterministic judging.

---

## Instruction

HACHATHON IS NOW LIVE! Feel free to access the instructions here: [INSTRUCTIONS](#)

---



## Prizes & Awards

\$10,000 in cash, three selective awards, and up to 50 paid opportunities.

### Prizes & Rewards



1st place reward

**\$5,000**



2nd place reward

**\$3,000**



3rd place reward

**\$2,000**

### Selective awards



#### Best Engineering Workflow

For the solution that stands out for how well it was built. Strong architecture, thoughtful technical decisions, clean implementation and a workflow that feels robust.



#### Most Useful Real-World Workflow

For the solution we could genuinely see people using. Ideas that solve a real problem in a practical, useful or unexpectedly creative way.



#### Best Demonstrated Improvement

For the clearest improvement from the original baseline. Make the progress obvious: what changed, why it changed, and what got measurably better.

## Opportunities at micro1

Up to 50 top performers may be considered for paid, flexible engineering work with micro1, subject to technical verification, eligibility and project availability. Additional strong performers may be invited to

complete micro1's accelerated AI interview. Every eligible participant with a valid submission receives a digital participation certificate.

If you're interested in taking a look at some of our openings, please see our most in demand engineering roles here:

[Senior Software Engineer](#)

[Open Source Contributor](#)

[Senior Back End Engineer](#)

Or browse our opportunities page [here](#).

## **Conditional coding-agent reimbursement**

After validation, micro1 may separately offer to acquire qualifying agent-use traces at \$2 to \$15 per trace, capped at \$100 to \$200 per participant. This is not part of the \$10,000 prize pool, does not affect judging, and is not guaranteed by registering or submitting. Any acquisition runs under separate terms covering artifacts acquired, permitted uses, rights granted, compensation, validation standard, privacy treatment and payment timing.

---

## **Stages & Timeline**

**Plan. Build. Prove. Submit.**

| Stage                          | Date and time (UTC)                                                                                   | What happens                                                                           |
|--------------------------------|-------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------|
| 1. Registration opens          |  Fri, Aug 21   15:00 | The HackerEarth page goes live and direct registration begins                          |
| 2. Pre-event briefing          |  Thu, Aug 27   15:00 | Rules, submission format, support process and logistics. No problem details            |
| 3. Hackathon Kickoff           |  Fri, Aug 28   15:00 | Event opening, you will receive the problem document and any starter material you need |
| 4. Final-day checkpoint        |  Sun, Aug 30   23:59 | Submission reminder, known issues and support escalation window                        |
| 5. Completion of the hackathon |  Mon, Aug 31   18:00 | Submissions close. Late or incomplete entries are not accepted                         |
| 6. Validation screen           |  Aug 31 - Sep 1      | Eligibility, completeness, reproducibility, plagiarism and trace-integrity checks      |
| 7. Judge review                |  Sep 2 - Sep 4       | Qualified submissions are scored with the published rubric                             |
| 8. Winners announced           |  Mon, Sep 7   17:00  | Winners, certificates and next-step notifications released                             |



Add to calendar



Kick off

Aug 28 | 15:00 UTC



Deadline

Mon, Aug 31 | 18:00 UTC

## Eligibility Criteria

1. Participants must be at least 18 years old at the time of registration.
2. Open globally except where participation or payment is prohibited by applicable law, sanctions, export controls, or organizer and platform restrictions.
3. Entries are individual only. A participant may register once and submit one final entry. Revisions are allowed until the deadline; only the latest complete submission is evaluated.
4. Participants should have at least six months of practical software-building experience or equivalent hands-on evidence. Professional employment is not required.
5. Eligible participants include working engineers, founders, open-source contributors, competitive programmers, final-year students, graduate students and recent graduates who meet the experience standard.
6. micro1 employees, event administrators, judges, anyone involved in creating or testing the challenge, and their immediate household members are not eligible for prizes.

- Participants must be able to receive payment through the approved payout rail in their country to receive a cash prize or trace-acquisition payment.
- Accurate identity, location, contact and eligibility information is required. Duplicate or false registrations may be disqualified.

## Evaluation Criteria

Scored out of 100 by micro1's engineering team.

**Qualification gate.** A submission is scored only after it passes eligibility, completeness, integrity, trace and reproducibility checks. A project that cannot be run or verified may be disqualified before rubric scoring.

| Criterion                    | Weight | What judges assess                                                                                                                          |
|------------------------------|--------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Problem & user value         | 15%    | Solves a meaningful problem for a clearly defined user. Who experiences the bottleneck, and why does solving it matter?                     |
| Agent solution & engineering | 30%    | Uses agents purposefully and is technically sound. Which design choices helped the agent solve the problem?                                 |
| End to end quality           | 20%    | Completes a realistic, self contained execution and produces a result the user can use. Would the intended user consider this high quality? |
| Measured improvement         | 15%    | Demonstrates gains over a fair baseline and uses the changelog to connect each iteration to evidence.                                       |
| Reproducibility              | 15%    | Gives another person a clear path to run the solution and baseline and reach the main result from a clean environment.                      |
| Hot take / insights          | 5%     | Turns an observed failure mode into a practical lesson for building more reliable agents.                                                   |

### Tie-break order

- Higher Agent Solution & Engineering score
- Higher Reproducibility score
- Higher Measured Improvement score
- Higher End to End Quality score
- Final panel review of documented evidence. The judges' decision is binding to the extent permitted by law and the official terms



---

## Submission Package

### Four items. All four are required.

1

Complete solution  
code & improvement  
changelog

2

Reproduction  
guide

3

Solution video,  
up to 5 minutes

4

Agent  
trajectories

#### Complete solution code and improvement changelog

Share the full project and everything required to run it, including the instructions that shape each agent. Use the README to introduce the intended user, explain their current bottleneck, and describe why solving it is valuable. Add a clearly labelled Improvement Changelog with an entry for every meaningful iteration, connected to the evidence that guided your next decision. Close with the main failure mode and your hot take.

#### Reproduction guide.

Written for someone starting from a clean environment. Walk them through setup and give the exact commands for the solution, the baseline and the evaluation. Explain which data is required and what output to expect. Share relevant versions along with approximate runtime and cost.

#### Solution video, up to 5 minutes.

Begin with the problem and the simple baseline, then walk through one realistic execution from start to finish. Show the final comparison and briefly explain the changelog. Highlight the change that contributed most, and one experiment you removed.

#### Agent trajectories.

Representative trajectories for every agent you used, easy to follow from the agent instructions through to the final result. Show what the agent did and how its tools responded. Capture the feedback that shaped its next step, plus any retries or human checkpoints.

---

# Rule Book

1. You are welcome to build with tools and components you already know.
2. Make it clear what existed before the competition and what you added.
3. Use every tool and component according to its licence and service terms.
4. Keep consequential actions controlled through a sandbox or simulation. Add human approval before the action happens.
5. Make a qualified human reviewer part of any solution that could significantly affect someone.
6. Choose a legal and ethical use case that treats people and their data responsibly.
7. Use information you are allowed to share. Public or synthetic data are usually easiest. Approved anonymous data also works.
8. Keep credentials and private information outside the submission.
9. Connect every claim about your results to the evidence you submit.
10. Give judges enough access to run the project and reproduce the main result.

***Submissions are governed by the Hackathon Participation Agreement (accepted at registration), under which micro1 owns submissions and may use them for AI model training and evaluation.***

---

## FAQs

**Q: Is registration free?**

A: Yes. There is no fee to register or submit.

**Q: Can I participate as a team?**

A: No. The August 2026 edition is for individual participation and individual prizes.

**Q: When will the exact problem be available?**

A: The exact problem will be revealed at kickoff on **August 28th**. Before then, micro1 will share the theme, rules, submission requirements, and schedule.

**Q: Are AI tools allowed?**

A: Yes. Coding-agent use is required. You must disclose the tools you use and submit the required trajectories or traces.

**Q: Do you provide API keys or model credits?**

A: No. Participants must use their own agent setup.



**Q: Which programming languages can I use?**

A: Any language is allowed, with **Python, TypeScript, Java, C++, Go, and Rust** recommended, subject to the runtime constraints in the final problem PDF.

**Q: Do I need to be looking for a job to participate?**

A: No. This is a competition. Any future interest in paid micro1 work will be handled separately.

**Q: What makes a submission valid?**

A: A valid submission must be **timely, complete, original, policy compliant, and reproducible**, and include the required repository, archive, tests, README, agent-use evidence, and demo video.

**Q: Does every submission receive a certificate?**

A: Yes. Every eligible participant with a valid submission receives a **digital participation certificate**.

**Q: Does submitting guarantee trace reimbursement or a job?**

A: No. Trace acquisition is a separate, limited programme if approved. Employment and project opportunities are separate and subject to additional review.

**Q: How much time will this take?**

A: Plan for a **focused three-day weekend**. Keep your problem scope tight so the deliverables are achievable within the timeframe.

**Q: Where can I get help?**

A: Use the **organizer contact and Help options** on the challenge page. Official clarifications will be shared with all participants.